

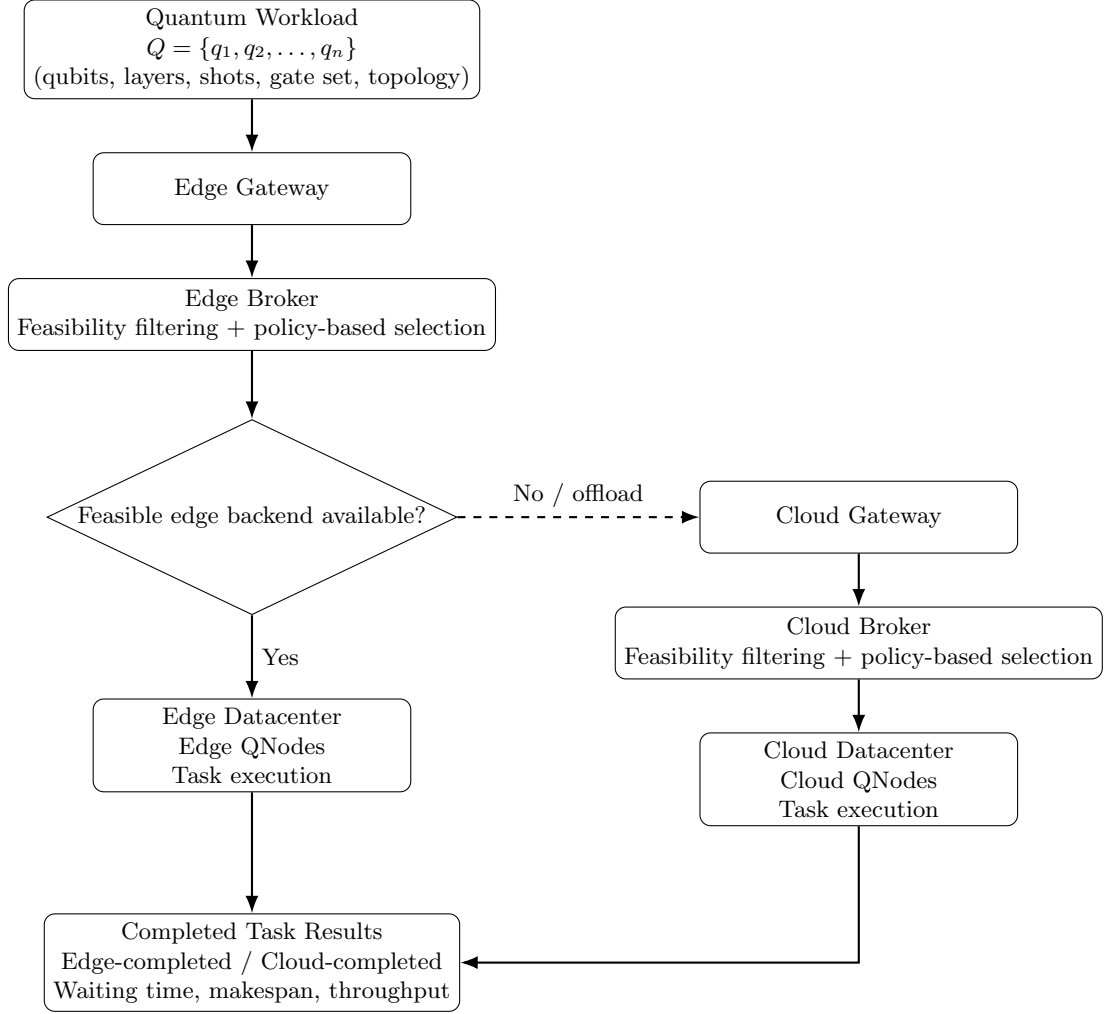


Figure 1: Broker-level scheduling workflow in the hybrid quantum edge-cloud framework. Tasks are first submitted through the edge layer, filtered for feasibility, and either assigned to edge backends or offloaded to the cloud when no feasible edge backend is available.

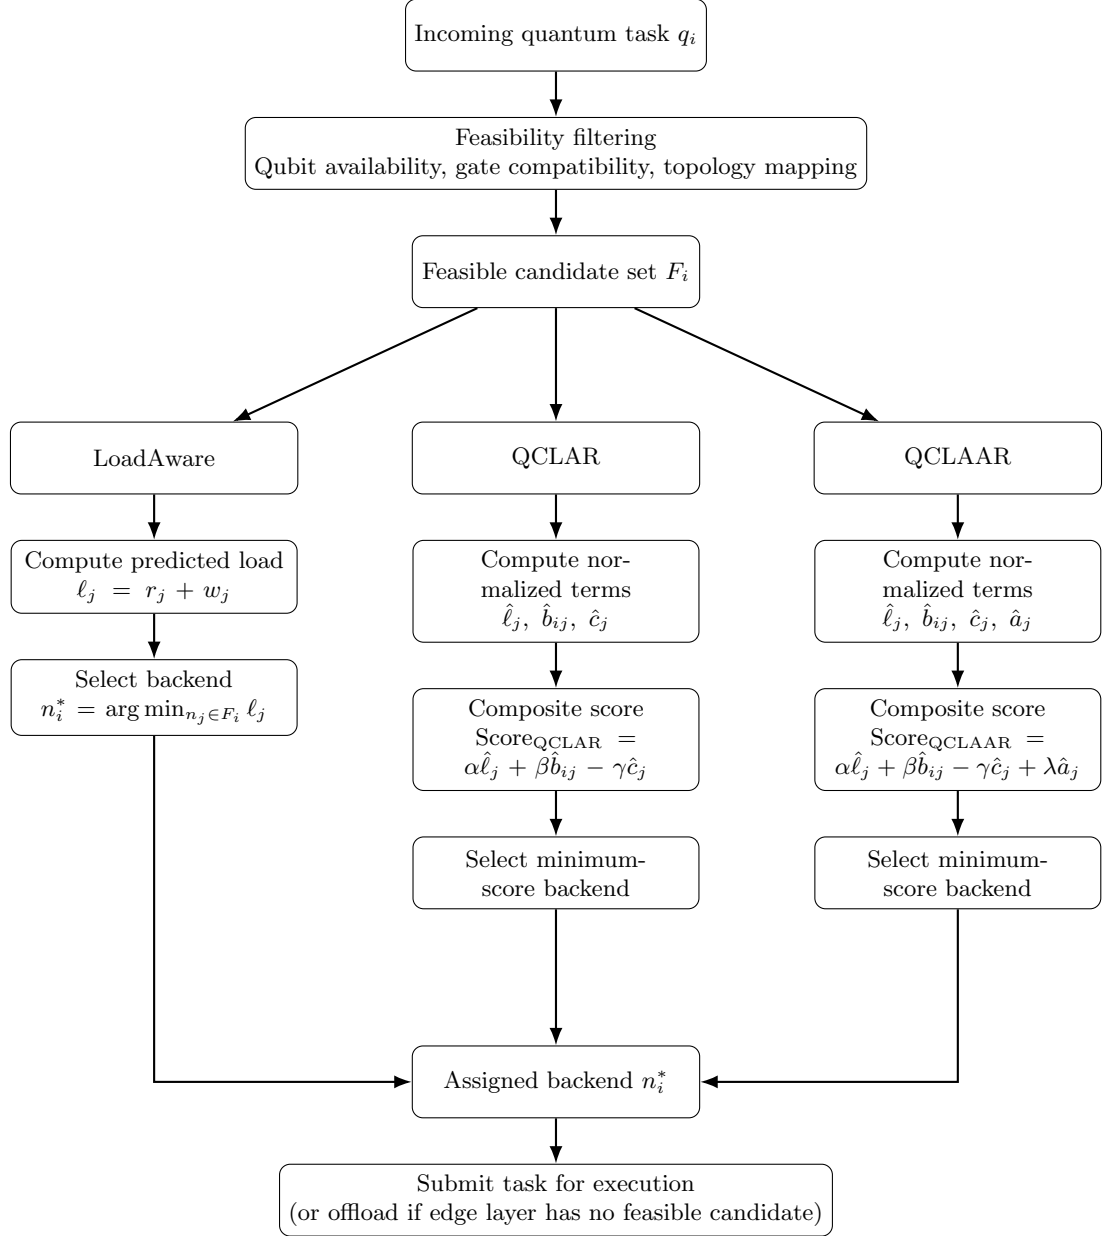


Figure 2: Unified decision pipeline for the compared broker policies. LoadAware selects the feasible backend with minimum predicted load, while QCLAR and QCLAAR extend this process with execution-burden, backend-capability, and queue-pressure-aware scoring.